DATA203 Foundational Python (Prof. Maull) / Fall 2026 / HW3

Points Possible	Due Date	Time Commitment (estimated)
20	Thursday, October 29 @ midnight	*up to* 15 hours

- **GRADING RUBRIC:** Grading will be aligned with the completeness of the objectives.

  Your solution is evaluated on:

  1. whether your solution correctly answers the question(s) being asked,
  2. using the Python language as required in the question(s).

  Your solution is **NOT** evaluated on:

  1. the elegant use of Python,
  2. optimal or high performance Python code unless *explicitly noted* in the question(s).

- **INDEPENDENT WORK:** Copying, cheating, plagiarism and academic dishonesty *are not tolerated* by University or course policy. Please see the syllabus for the full departmental and University statement on the academic code of honor.

- **LARGE LANGUAGE MODEL (LLM) / ARTIFICIAL INTELLIGENCE/LLM (AI) POLICY:**

  LLMs / Artificial Intelligence technologies such as Claude, Gemini, Grok, etc are **NOT** to be used in the production of your final homework solutions.

  Such technologies offer important capabilities for your *learning* journey, but they are not appropriate in this course for the production of your solutions.

  You *may* use such technologies to *learn* more about specific Python concepts or to help clarify your knowledge about the details of Python syntax, modules, techniques or other programming principles.

  Final solutions found to be the *partial or entire product* of LLM/AI use **will receive 0 points**, without clarification, resubmission or correction.

OBJECTIVES

- PART 1: Convert unstructured data to structured data.
- PART 2: Use the data to produce rating of chocolates by country.

WHAT TO TURN IN

You will enjoy the highest benefits of the starter notebook if you clone the HW Github repository from your Jupyter Hub terminal with the command:

```
git clone https://github.com/kmhuads/f26_data203.git
```

This will ensure you have the most updated files and starter notebook.

Once you have cloned the repository, you can edit the starter notebook with your solution code.

When you are done with your work, it will be best to zip your `hw2` folder and all sub-folders with the terminal command (one level outside your notebook folder):

```
zip -r data203_hw3_maull.zip ./hw3
```

This will produce the file with all necessary supporting files (notebooks, data output, etc.) then download it from the Jupyter Hub to your local machine, then upload the `.zip` to Teams.

If you are confused on how to do this, please ask, or visit one of the many tutorials on the basics of using zip in Linux.

If you choose not to use the provided notebook, you will still need to turn in a `.ipynb` Jupyter Notebook and corresponding files according to the instructions in this homework.

ASSIGNMENT TASKS

(75%) PART 1: Convert unstructured data to structured data.

ASSIGNMENT BACKGROUND

In our previous assignments we learned about YAML and JSON files for storing data.

We did not explicitly say what kind of data was in the file, but instead we used Python objects like lists and dictionaries to easily store the data into formats which were compatible and interchangeable. The Python objects dictated what was in the data. We could, for example, load a JSON file into a Python dictionary or save a Python dictionary into a JSON file.

The dictionary is what is called a *structured* format – that is to say it's components can be easily identified (or inferred) and more importantly, that it is easily consumable by a machine (your Python code).

You might ask, if these are structure then what is *unstructured*? Unstructured data differs fundamentally in that first you will not easily understand the format or components of the data without deeper prior knowledge. For example, most written documents (letters, manuals, publications) are *unstructured*. They are not easily understood without prerequisite knowledge – for example, if you do not know Spanish, you will have a hard time reading Spanish words or even understanding the meaning, structure and intent of the sentences.

Machines are equally clueless with unstructured data, except when we use very sophisticated software such as the LLMs (Large Language Models) in ChatGPT, Claude, Gemini and others, but know that getting a machine to process this unstructured data in the way current generation LLMs do, is a very heavy lift and requires tremendous amounts of resources to accomplish.

You might also wonder why unstructured data is such a big deal or why we are concerned about it here. Take a look at some remarkable statistics about the state of structured and unstructured data in today's data landscape. According to IDC's (International Data Corp) 2024 report *Worldwide Global DataSphere Structured and Unstructured Data Forecast, 2024–2028*:

1. By the end of 2026, unstructured data is expected to account for nearly 88% of global data generation.
2. Unstructured data CAGR (compound annual growth rate) is projected to be 21.4% in the five year period between 2023-2028, topping nearly 400 zettabytes (ZB) of data.
3. By 2028, structured data is still only expected to make up less than 20% of all data generated.

With statistics like these – a well positioned data scientist is one who at least understands the landscape, processes and tools of structured and unstructured data. As a side note, you might be scratching your head on how much data a "zettabyte" is. Your laptop computer may have a total of 512GB of storage on its hard drive. A zettabyte is 2 *billion* of such laptop hard drives **filled**.

DATASET

We are going to warm up first with some fun data which will please the chocolate lovers among us (yes, including myself with a strong preference for dark, >70% cacao). A few years ago, I found a wonderful website called *Flavors of Cacao* (FoC) and on that website is a page of over 2800 chocolate bars which have been taste tested and rated by the Manhattan Chocolate Society. Visit the full website and ratings page:

- Flavors of Cacao - main site
- FoC / MCS Chocolate Bar Database and Ratings

You may notice something curious, though. The database ratings page is just one big web table and for all intents and purposes is difficult to filter, sort or search – unless you just want to CTRL-F your way to a slow time hunting. This is a shame, but also an amazing learning opportunity as this is a *prime* example of unstructured data – yes, you could argue it is semi-structured since it is *presented* as a table, but it is otherwise not easily consumed by a machine. We are going to convert this page into something more useful – a CSV file.

§ Task: 1.0 Convert the *FoC* table to a CSV file.

You have, no doubt, come across CSV files at least once in your lifetime so far. CSV stands for "Comma Separated Values" and like you probably remember, this file usually contains a bunch of data like this example:

```
height (cm),weight (kg),age
165,60,30
```

```
175,75,45
180,80,50
155,55,25
170,65,35
```

You probably used Excel to load the file and went about the task of doing something useful with it. This is an example of structured data. You also have hints and clues about what each number means and luckily the first line of the file includes the units of those numbers (cm, kg and we'll assume age is in years). Each column is organized specifically around a data variable and each row contains a single data point across those variables. CSV files can be large and are valuable for storing many types of data that can be organized in a table.

We are going to convert the chocolate bars table on the FoC page to a CSV and to do so we are going to need to import a very special library called BeautifulSoup to do it. Please load the starter notebook for this assignment to learn how to use it.

Your task is to convert the FoC ratings table to a CSV file and store the data in a file within the `data/` folder called `foc_chocodb.csv`.

Here are the specifics:

1. The file will contain 10 columns:

- REF
- Company (Manufacturer)
- Company Location
- Review Date
- Country of Bean Origin
- Specific Bean Origin or Bar Name
- Cocoa Percent
- Ingredients
- Most Memorable Characteristics
- Rating

2. The file will use comma `,` as the column (field) delimiter.
3. All fields (columns) in the final output will be surrounded by the quote character `"`. If there is a `"` character in a field, replace it with two quotes or `""`. For example, if you have something like `Fine "Cacao" Corp` in a field, your output will be `"Fine ""Cacao"" Corp"`.
4. The file will have 2842 rows of data represented by the entire table.
5. The file will be stored in the `data/` folder of your notebook homework and named `foc_chocodb.csv`.

§ Task: 1.1 Create a CSV file with just the chocolates with `Rating` equal to or higher than 3.9.

You will loop over all the data and check just the `Rating` and store only those rows where the `Rating` is greater than or equal to 3.9. Name the file `data/foc_chocodb_top_rated.csv`.

You will find and use the Python `csv` module useful! Read the documentation about it here. You may **not** import any other libraries in this part:

- Python `csv` module documentation

You may use the `csv.reader()` function, just remember that the `Country of Bean Origin` is index 4 and `Rating` is index 9 (or -1). You can also use `csv.DictReader()` but you will still need to store the data in the CSV file.

§ Task: 1.2 Write a function that reads the file and returns the rows of countries as a string given the input parameters.

Your function will take up to five parameters, a filename *fname*, a minimum rating *rating_min*, a maximum rating *rating_max* and a header indicator *header*.

Your code will open the file, scan the `Country of Bean Origin` column for the *country* string and/or `Rating` column for rating and return the lines that match the values of the input parameters.

You should normalize the country using lowercase in both the file and the input.

FUNCTION SPECIFICATION

function **NAME**: `filter_db_csv`

function **INPUT**:

- *fname* → the string for the filename containing the CSV Data
- *country* → the string containing the country name to filter on
- *rating_min* (optional; default 0) → a float containing the *minimum* rating
- *rating_max* (optional; default 4) → a float containing the *maximum* rating
- *header* (optional; default `True`) → `True` if the output is expected to include the header, otherwise `False`

function **OUTPUT**:

- a string which contains the header (first line of the CSV file With the column names)
- all lines in the CSV input file *fname* which match the parameters given
- remember the input *country* uses the `Country of Bean Origin` column
- the input *rating_min* and *rating_max* use the `Rating` column
- each line of data must have a terminating newline `\n`

Code Example:

```
filter_db_csv_str("foc_chocodb.csv", country="uganda")
```

Example Output:

NOTE: the output has been reduced so it can be more readable here, but you will need to emit *all* columns and data.

```
"REF","Company (Manufacturer)",..,"Country of Bean Origin","Specific Bean Origin or Bar Name",...,"Rati
"2546","5150",...,"Uganda","Semuliki Forest, batch 1",...,"3.25"
"2832","Bonnat",...,"Uganda","Uganda, Geant des Brames",...,"3.25"
...
```

§ Task: 1.3 Use your function to produce a file with just the parameters given.

1. Produce a new file `data/foc_chocodb_pancubrabel.csv` that includes **only** the following countries:

- Panama
- Cuba
- Brazil
- Belize

2. Produce a new file with only those chocolates with `Rating` less than 2. Name the file `data/foc_chocodb_low_tier.csv`.
3. Produce a new file with only those chocolates from *Brazil* with a `Rating` greater than or equal to 3.5. Name the file `data/foc_chocodb_brazil_high_tier.csv`.

Don't overthink it – remember that you can use `write()` to write data to a file, given a string.

§ Task: 1.4 Use your function to filter on the same input but output the data list of dictionaries.

You will do the same as before, but the return result will be a list of dictionaries.

You may use the Python `csv.DictReader()` function to save a tremendous amount of time.

FUNCTION SPECIFICATION

function **NAME**: `filter_db_dict`

function **INPUT**:

- *fname* → the string for the filename containing the CSV Data
- *country* → the string containing the country name to filter on
- *rating_min* (optional; default 0) → a float containing the *minimum* rating
- *rating_max* (optional; default 4) → a float containing the *maximum* rating

function **OUTPUT**:

- a list of dictionaries
- each entry in the list is a dictionary which contains the header (first line of the CSV file With the column names) as the key and the corresponding value
- all outputs come from the input file *fname* which match the parameters given
- remember the input *country* uses the `Country of Bean Origin` column
- the input *rating_min* and *rating_max* use the `Rating` column

Code Example:

```
filter_db_dict("foc_chocodb.csv", country="uganda")
```

Example Output:

```
[{'REF': '2546',
'Company (Manufacturer)': '5150',
'Company Location': 'U.S.A.',
'Review Date': '2021',
'Country of Bean Origin': 'Uganda',
'Specific Bean Origin or Bar Name': 'Semuliki Forest, batch 1',
'Cocoa Percent': '80%',
'Ingredients': '3- B,S,C',
'Most Memorable Characteristics': 'mildly bitter, basic cocoa, fatty',
'Rating': '3.25'},
...
]
```

(25%) PART 2: Use the data to produce rating of chocolates by country.

We are almost there. What's left is to combine our newly acquired knowledge about chocolates to interactive visualizations.

§ Task: 2.0 Create a dictionary of the average chocolate ratings by country.

Use your previous function `filter_db_dict()` to produce the dataset for the following 20 countries:

- Congo, Jamaica, Panama, India, Indonesia, Liberia, Honduras, Sao Tome & Principe, Costa Rica, Tanzania, Dominican Republic, Malaysia, Papua New Guinea, Haiti, China, Nicaragua, Peru, Taiwan, Congo, India

You will output a dictionary of the form:

```
{
  'Haiti': 3.217,
  'Gabon': 3.25,
  ...
}
```

You may write a function for this or you can run it in a cell as long as you store the output so it can be used in the next step.

§ Task: 2.1 Use the dictionary from Task 2.0 to build an interactive map.

We will use a visualization tool called Folium to display the countries, ratings and a color coded marker.

First, take a look at Folium and see how cool it is:

- Folium Python module

You will need to install Folium by running:

```
!pip install folium
```

Your map will look something like:



To produce this map, look at the scaffolding code in the starter notebook.

You will need to load the file `data/country_locations.csv` and use the latitude and longitude data to place the markers correctly on the map. To do this, you may want to load it into a structure with a key on the country, so you can get the latitude and longitude.

Here are the requirements for your map:

1. An average rating under 3 should be colored *red*;
2. color the markers *orange* when the average rating is between 2.99 and 3.14;
3. color the markers *green* when the average rating is between 3.15 and 3.23;
4. color the markers *purple* when the average rating is greater than 3.24.